

Input/Output Mechanics

Based on slides © McGraw-Hill
Additional material © 2004/2005 Lewis/Martin
Modified by Diana Palsetia

Examples of Input/Output (I/O) Devices

Standard Interfaces

1. User output
 - Console or Prompt (just text), Video Display, printer, speakers
2. User input
 - Keyboard, mouse, trackball, game controller, scanner, microphone, touch screens, camera (still and video), game controllers
3. Storage
 - Disk drives, CD & DVD drives, flash-based storage, tape drive

More...

Communication

- Network (wired, wireless, optical, infrared), modem

Sensor inputs

- Temperature, vibration, motion, acceleration, GPS
- Barcode scanner, magnetic strip reader, RFID reader

Control outputs

- Motors, actuators

CIT 593

2

How to Perform I/O?

Require specialized knowledge and protection

- **Knowledge** of I/O device work
 - programmers don't want to know this!
 - unless you programmer who is writing device drivers
- **Protection** for shared I/O resources (e.g. Hard- Disk) want process isolation

Solution: **service routines** or **system calls**

- Low-level, privileged operations performed by operating system
- Used by almost all applications and hence made a routine
- Example: keyboard & display polling routine for LC-3
- Example: getc, puts in C

CIT 593

3

LC-3 I/O

I/O is done via the TRAP instruction

Code	Equivalent	Description
OUT	TRAP x21	Write one character (in R0[7:0]) to console.
GETC	TRAP x20	Read one character from keyboard. Character stored in R0[7:0].
IN	TRAP x23	Read (and echo) one character from keyboard. Character stored in R0[7:0].
PUTS	TRAP x22	Write null-terminated string to console. Starting address of string is in R0.

CIT 593

4

EXAMPLE

;Prompt the user to enter a number, quits when pressed 3, prints DONE before terminating

```
.ORIG x3000
LD R1, VAL3
NOT R1, R1
ADD R1, R1, #1
LEA R0, ENTER      ;R0 contains address of first char of string
PUTS                ; displays the string
LOOP: IN            ; read and echo (R0 contains the char)
      AND R3, R3, #0
      ADD R3, R1, R0
      BRnp LOOP
      LEA R0, ENDMSG ; R0 = starting address of string
      PUTS
      HALT
ENTER: .STRINGZ "Enter a number from 1 to 9:\n"
ENDMSG: .STRINGZ "\nDONE"
VAL3:  .FILL x0033
.END
```

CIT 593

5

System Call

1. User program invokes system call i.e. calls the part of the system code (O.S code) that deals with I/O
 - E.g. TRAP x20 (GETC) in LC3
2. Operating system code performs operation
 - TRAP routine follows calling conventions (callee-save) just like regular user routine
3. Returns control to user program

CIT 593

6

LC3 TRAP Instruction

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
TRAP	1	1	1	1	0	0	0	0	trapvect8							

Trap vector

- Identifies which system call to invoke
- Serves as index into table of service routine addresses
 - LC-3: table stored in memory at **0x0000 – 0x00FF**
 - 8-bit trap vector zero-extended to form 16-bit address, thus allows upto 256 service routines

Where to go

- Lookup starting address from table and place in it PC

Enabling return

- Save address of next instruction (current PC) in R7

How to return

- Place address in R7 in PC

CIT 593

7

LC3 Trap Vector Table

Trap Vector Table

- Used to associate code with trap number
 - The table contains the address of where the service routine is located
- The table is stored in memory (**x0000** through **x00FF**)

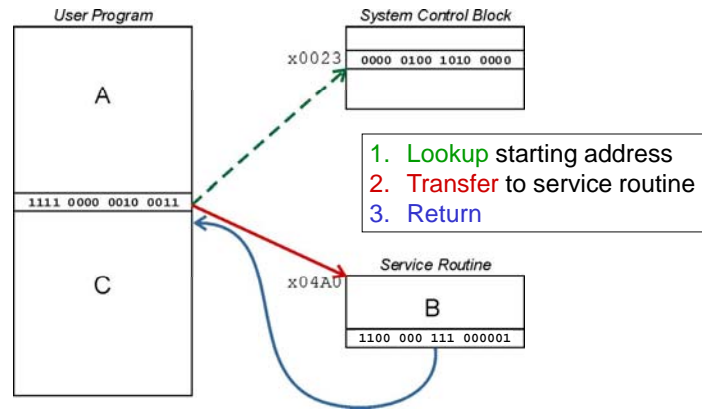
Vector Table (Memory location)	Routine location	Routine
x21	x0430	output a character to the monitor
x23	x04A0	input a character from the keyboard
x25	xFD70	halt the program (HALT)

Note: The service routines are located at location at x0200 - x2FFF in memory except the HALT routine

CIT 593

8

TRAP Mechanism Operation



CIT 593

9


I/O Basics

Control/Status Registers

- CPU tells device what to do -- write to control register
- CPU checks whether task is done -- read status register

Data Registers

- CPU transfers data to/from device



Device electronics

- Performs actual operation
 - Pixels to screen, bits to/from disk, characters from keyboard

CIT 593

10

Programming Interface

How are device registers identified?

- Memory-mapped vs. special instructions

How is timing of transfer managed?

- Asynchronous vs. synchronous

Who controls transfer?

- CPU (polling) vs. device (interrupts)

CIT 593

11

Memory-Mapped vs. I/O Instructions


Instructions

- Designate opcode(s) for I/O
- Register and operation encoded in instruction



Memory-mapped

- Assign a memory address to each device register
- Use data movement instructions (LD/ST) for control and data transfer
- Hardware intercepts these address
- No actual memory access performed



CIT 593

12

Transfer Timing

Synchronous

- Data supplied at a fixed, predictable rate
- CPU reads/writes every X time units
- Infrequently used because of speed difference
 - I/O is much **slower** than the processor
 - E.g. A 300 Mhz processor can execute an instruction every 33 nanoseconds
 - E.g. Humans don't type as fast as a processor can read

Asynchronous

- Data rate less predictable
- CPU must synchronize with device, so that it doesn't miss data or write too quickly
- Handles the speed mismatch

CIT 593

13

Transfer Control

Who determines when the next data transfer occurs?

Polling

- CPU keeps checking status register until new data arrives OR device ready for next data
- “Are we there yet? Are we there yet? Are we there yet?”

Interrupts

- Device sends a special signal to CPU when new data arrives OR device ready for next data
- CPU can be performing other tasks instead of polling device
- “Wake me when we get there.”

CIT 593

14

LC-3 I/O

Memory-mapped I/O (Table A.3 in Appendix A of text)

Location	I/O Register	Function
xFE00	Keyboard Status Reg (KBSR)	Bit [15] is 1 when keyboard has received a new character.
xFE02	Keyboard Data Reg (KBDR)	Bits [7:0] contain the last character typed on keyboard.
xFE04	Display Status Register (DSR)	Bit [15] is 1 when device ready to display another char on screen.
xFE06	Display Data Register (DDR)	Character written to bits [7:0] will be displayed on screen.

Asynchronous devices

- Synchronized through status registers
 - Two ways: **Polling** and **Interrupts**
- We'll talk first about polling, a bit on interrupts later


CIT 593

15

Input from Keyboard

When a character is typed:

- Its ASCII code is placed in bits [7:0] of KBDR (bits [15:8] are always zero)
- The “ready bit” (KBSR[15]) is set to 1
- Keyboard is disabled -- any typed characters will be ignored



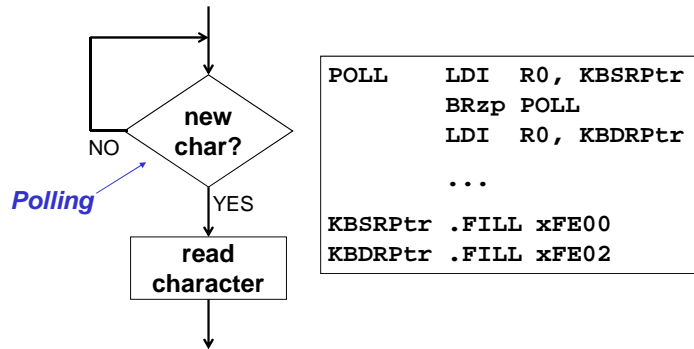
When KBDR is read by processor:

- KBSR[15] is set to zero
- Keyboard is enabled

CIT 593

16

Basic Input Routine (GETC)



CIT 593

17

General I/O

The interactions described can be applied to most I/O devices

- The status register in a printer might tell you (via computer)
 - If the printer is printing, or
 - Is out of paper, or out of toner

Data/Information that is transferred can be more than just a byte (known as Block I/O)

CIT 593

18

Role of the Operating System

How does non-privileged code perform I/O?

- Answer: it doesn't; it asks the OS to perform I/O on its behalf

How is this done?

- Making a system call into the operating system

In real systems, only the operating system (OS) does I/O

- "Normal" programs ask the OS to perform I/O on its behalf

Hardware prevents non-operating system code from

- Accessing I/O registers
- Operating system code and data
- Accessing the code and data of other programs

Why?

- Protect programs from themselves
- Protect programs from each other
- Multi-user environments

CIT 593

19

Memory Protection

Code in **privileged/supervisor** mode

- Can do *anything*
- Used exclusively by the operating system

Code in user mode

- Can't access I/O parts of memory
- Can only access some parts of memory

Division of labor

- OS - make policy choices
- Hardware - enforce the OS's policy

CIT 593

20

OS and Hardware Cooperate for Protection

Hardware support for protected memory

- Consider a n-bit protection register (MPR: memory protection register)
 - Each bit protects some region of memory

When a processor performs a load or store

- Checks the corresponding bit in MPR
- If MPR bit is not set (and not in privileged mode)
 - Trigger illegal access


The OS must set these bits before running each program

CIT 593

21

MPR in LC3 Example

- Memory access is limited by memory protection register (MPR)
- Each MPR bit corresponds to 4K memory segment
- 1 indicates that users can access memory in this segment




CIT 593

22

Privilege

Privilege: Processor modes

- Privileged (supervisor) vs. Unprivileged (user) indicated by bit
- E.g. in LC3
 - Encoded in 15th bit of Processor Status Register (PSR)
 - PSR has other info such as Condition Codes (NZP) and Priority Level(PL)



CIT 593

23

Managing Privilege

Who sets privilege bit in PSR?

- TRAP instruction

Who clears privilege bit?

- An instruction designed to return from trap routines
- Is different from **user subroutine return instruction** as a user subroutine

CIT 593

24

Interrupt-Driven I/O

Timing of I/O controlled by *device*

- Tells processor when something interesting happens
 - Example: when character is entered on keyboard
 - Example: when monitor is ready for next character
- Processor *interrupts* its normal instruction processing and executes a service routine (like a TRAP)
 - Figure out what device is causing the interrupt
 - Execute routine to deal with event
 - Resume execution
- No need for processor to poll device
 - Can perform other useful work

CIT 593

25

To implement an Interrupt mechanism


1. Way for I/O device to **signal** CPU that its wants service

- We already have a **Ready** bit indicating that



2. Way to know that device **has a right to request service**

- **Interrupt Enable (IE)** bit in device register is set by processor
 - Depending on whether processor wants to give service
- When ready bit is set and IE bit is set, interrupt is signaled
I/O device will get service



CIT 593

26

To implement an Interrupt mechanism (contd..)

3. Whether I/O device has higher **priority** among multiple I/O requests AND with the current program in execution

- Every instruction executes at a stated level of urgency
- LC-3: 8 priority levels (PL0-PL7) with 7 being higher priority
- Example:
 - current program runs at PL0
 - Nuclear power correction program runs at PL6
 - It's OK for PL6 to interrupt PL0 program, but not the other way around
- A special hardware compare priority levels of the interrupts generated, the one with higher priority gets to use the processor
- The priority level is encoded in the PSR (Processor Status Register)

CIT 593

27

Maintaining the state of the machine

You were in the middle of adding 2 numbers in your program, but there was interrupt which needs to get serviced

- Did I complete my operation?
- After I resume, is my data in the registers that I was storing the same as I left of ?

Before servicing the interrupt

- Save state of the machine, i.e. Registers, PC, mode and CCs
 - This way I can come back start executing where I left

To service interrupt:

- Load the PC with starting address of the program that is to carry out the requirements of the I/O

CIT 593

28


Interrupt Implementation using LC3 Example

Interrupt vector: INTV

- 8-bit identification for device
- Also used as index into Interrupt Vector Table, which gives starting address of *Interrupt Service Routine (ISR)*
 - Just like Trap Vector Table and Trap Service Routine
 - In LC3 Interrupt Vector Table (IVT): x0100 to x01FF

What if more than one device wants to interrupt?

- External logic controls which one gets to go




CIT 593

29

How Does Processor Handle an Interrupt?

After completing current instruction, If INT (interrupt) =1, then

- Don't fetch next instruction
- Save state (PC, PSR (privilege and CCs) etc) on *Supervisor stack*
- Update PSR (set privilege bit i.e. PSR[15] = 1)



- Index INTV into IVT to get start address of ISR (put in PC)

After service routine

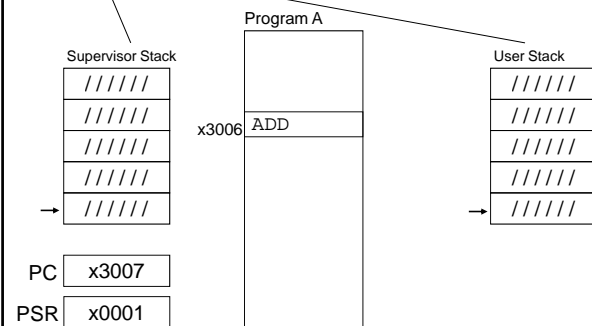
- RTI (Return from Interrupt) instruction restores machine state from *Supervisor stack*
 - Not the user stack to protection & isolation from users accessing operating system data

CIT 593

30

Example (1)

Different sections in memory

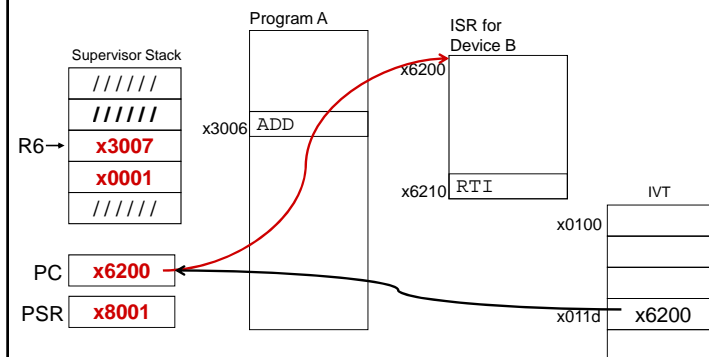


Executing ADD at location x3006 when Device B interrupts (INTV = x1d)

CIT 593

31

Example (2)

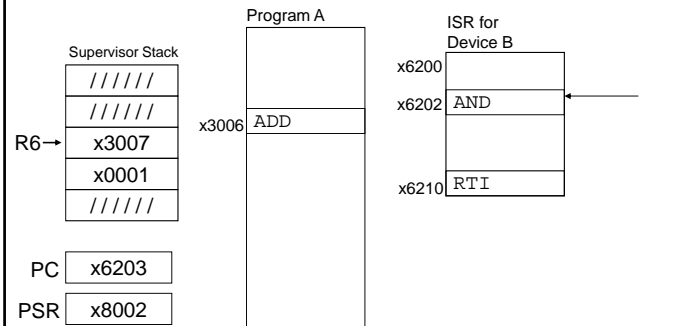


Push PC and PSR onto stack, set privilege bit, find Device B (INTV=x1d) service routine address in IVT, and transfer control
Note: Have not shown other registers on supervisor stack

CIT 593

32

Example (3)

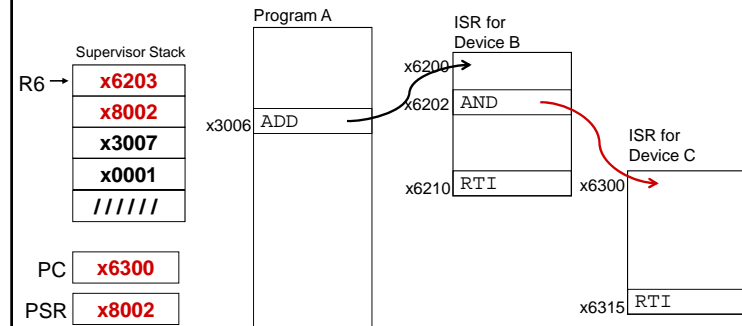


Executing AND instruction of Interrupt Routine for device B at x6202 when Device C interrupts

CIT 593

33

Example (4)

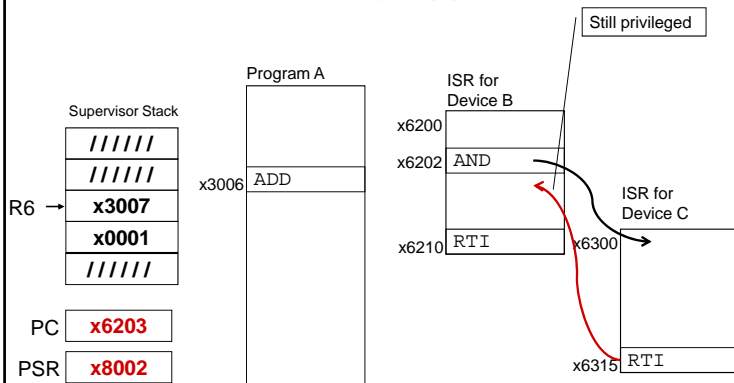


Push PC and PSR onto stack, set privilege bit, then transfer to Device C service routine (at x6300)

CIT 593

34

Example (5)

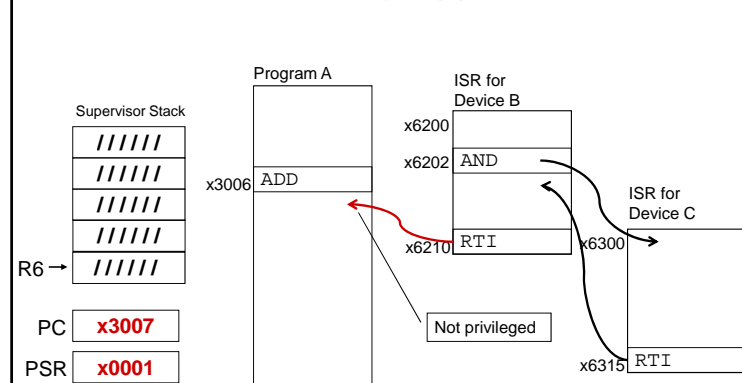


Execute RTI at x6315; pop PSR and PC from stack

CIT 593

35

Example (6)



Execute RTI at x6210; pop PSR and PC from stack; continue Program A as if nothing happened!

CIT 593

36

Software Interrupts

So far we talking about hardware interrupts i.e. interrupts from I/O devices.

A software interrupt a.k.a *exception* is interrupt caused when something unusual happens inside processor

- E.g. Divide by zero, overflow etc.